

DE 414: Practical 4 Report

Amy McDermott (26911264)

April 27, 2026

Introduction

This report documents the implementation of a Multi-layer Perceptron using Python. It explores the manual application of the forward pass, backpropagation, and loss calculation. The model's performance and ability to generalise are tested against the datasets `simple` and `two_blob`. By visualising the decision surfaces generated by the model, we can observe how the neural network's architectural choices - such as the number of hidden layers and neurons - impacts the network's capacity to handle various levels of data complexity.

Question 1: Single layer perceptron

In this section, a single-layer perceptron was implemented and trained on the `simple` dataset. This architecture consists of 2 input units, $P1 = 3$, and 1 output unit, $Q1 = 1$.

Loading of data

The first step of the practical involved loading the `simple` dataset and visualising its distribution. Using the `matplotlib` library, a scatter plot, shown in Figure 1, was generated with points colour-coded by their respective labels (red for class 0 and green for class 1).

Upon inspection of the resulting plot, the dataset consists of two clearly separated clusters. Due to the fact that a straight line can be drawn to perfectly divide the two classes without any overlap, this dataset is classified as linearly separable. On account of this, the single-layer perceptron (using a linear transformation and sigmoid activation) implemented in this section, is sufficient to achieve optimal classification performance.

Forward pass

The forward pass is the fundamental process by which the network transforms input data into a prediction. For the single layer implemented in this section, a linear layer output, z_i , and then the activation function output v_i is calculated. As demonstrated in Listing 1 below, z_i is computed by performing a matrix multiplication between the transpose of the weight matrix W and the input vector u_i (including a bias term). Following the linear transformation, the sigmoid activation function is applied to z_i to produce the final activation output, v_i .




Figure 1: Scatter plot of the `simple` dataset, demonstrating linear separability.

```

1 z_i = W.T @ u_i # Linear layer output (Transpose W to match dimensions
    for matrix multiplication)
2 v_i = sigmoid(z_i) # Activation function output
3 return z_i, v_i

```

Listing 1: Implementation of the forward pass for a single layer

Loss derivative

The optimisation of the network relies on the calculation of the gradient of the loss function (Binary Cross-Entropy) with respect to the model's predictions $\hat{y}$. This derivative serves as the initial error signal, λ , that is propagated backward through the network. As shown in the Listing 2 implementation below, the derivative of the binary cross-entropy loss is calculated using the following formula.

$$\frac{\partial L}{\partial \hat{y}} = \frac{y - \hat{y}}{\hat{y}(1 - \hat{y}) + \epsilon} \quad (1)$$

To ensure numerical stability and prevent a division-by-zero error when $\hat{y}$ is exactly 0 or 1, a small constant ϵ (1×10^{-7}) is added to the denominator.

```

1 eps = 1e-7
2 dL_dy_hat = None
3 dL_dy_hat = (y - y_hat) / (y_hat * (1 - y_hat) + eps) # Add eps to avoid
    division by zero
4 return dL_dy_hat

```

Listing 2: Implementation of the derivative of the binary cross-entropy

Backward pass

The backward pass is the mechanism by which the network propagates the error signal from the output layer back toward the input, enabling the update of the weights via gradient descent. This process begins by computing the gradient of the loss with respect to the linear output z_i , where Λ is the error signal.

$$\delta_i = \Lambda \odot v_i(1 - v_i) \quad (2)$$

This result is subsequently used to determine the gradient of the loss with respect to the weights, $\frac{\partial J}{\partial W_i}$.

$$\frac{\partial J}{\partial W_i} = u_i \delta_i \quad (3)$$

A new Λ is then calculated to be passed to the preceeding layer. However this is only necessary for deeper architectures, where more than one layer makes up the network. When this updated Λ is computed, the first row of the weight matrix W_i is excluded since this row exists to incorporate a bias term and does not receive an error signal from previous layers. Listing 3 shows the necessary code to execute the backpass functionality.

```
1 delta = Lambda * v_i * (1 - v_i) # Backpropagation factor for layer i (  
   Q_i,1)  
2 dJ_dWi = u_i @ delta.T # Gradients of weights for layer i  
3 Lambda = (W_i[1:, :] @ delta).flatten() # Backpropagation factor for  
   previous layer (Q_i,1)  
4 return dJ_dWi, Lambda
```

Listing 3: Implementation of the backward pass

Model training

Following the implementation of the forward and backward passes, the model was trained for 100 epochs using the default hyperparameters (learning rate of 0.1 and batch size of 1). Based on the training output, the model performance evolved as follows:

1. Epoch 0: test loss of -0.1529
2. Epoch 99: test loss of -0.0205

Table 1 illustrates model output over the last 10 epochs of training.

Table 1: Training and test loss for the final 10 epochs of the single-layer perceptron on the `simple` dataset.

Epoch	Training Loss	Test Loss
90	-0.0363	-0.0204
91	-0.0362	-0.0200
92	-0.0361	-0.0208
93	-0.0359	-0.0218
94	-0.0358	-0.0220
95	-0.0357	-0.0214
96	-0.0355	-0.0202
97	-0.0352	-0.0195
98	-0.0352	-0.0202
99	-0.0351	-0.0205

In this specific implementation, a value closer to zero indicates higher likelihood and therefore a better model fit. It is evident that the significant reduction in loss over model training demonstrates that the single-layer perceptron successfully converged.

Model decision surface

After training the model, the decision surface was visualised. The resulting plots, shown in Figure 2 displays a clear, linear separation between the two classes, which confirms that the single-layer perceptron has successfully learned to partition the data. It is important to note that the decision




Figure 2: Decision surface of the `simple` data

surface of a single-layer perceptron is linear. More complex decision boundaries can only be captured using multi-layer architectures. Considering the distribution of the `simple` dataset, as seen in

Figure 1, a linear decision surface, and hence a single-layer perceptron, is sufficient to achieve optimal classification performance.

Question 2: Multi-layer perceptron

Data loading and visualisation

The `two_blob` dataset was loaded and visualised using the same approach as the `simple` dataset. The resulting scatter plot, shown in Figure 3, reveals a more complex distribution of data points, with two distinct clusters that are not linearly separable. This non-linear distribution necessitates the use of a multi-layer perceptron, which can capture more complex non-linear decision boundaries than a single-layer architecture.




Figure 3: Scatter plot of the `two_blob` dataset, demonstrating non-linear separability.

Model architecture

To address the non-linear nature of the `two_blob` dataset, a multi-layer perceptron architecture was implemented. This architecture includes two extra layers W_2 and W_3 with dimensions $P_2 = 8$, $P_3 = 8$, and $Q_3 = 1$ respectively. The weights are then initialised for all three layers. Listing 4 below shows the code used to define the dimensions of the multi-layer architecture.

```
1 P_1 = 2 # Input layer has two features
2 Q_1 = 8 # Number of neurons in first hidden layer = number of outputs
3
4 P_2 = 8 # Number of inputs to second hidden layer = number of neurons in
   first hidden layer
5 Q_2 = 8 # Number of neurons in second hidden layer = number of outputs
6
```

```

7 P_3 = 8 # Number of inputs to output layer = number of neurons in second
  hidden layer
8 Q_3 = 1 # Final output for binary classification
9
10 # Initialise weights for all three layers
11 W_1 = np.random.uniform(-1/np.sqrt(Q_1),1/np.sqrt(Q_1),size=(P_1+1,Q_1)) #
  initialise weights for layer W_1
12 W_2 = np.random.uniform(-1/np.sqrt(Q_2),1/np.sqrt(Q_2),size=(P_2+1,Q_2)) #
  initialise weights for layer W_2
13 W_3 = np.random.uniform(-1/np.sqrt(Q_3),1/np.sqrt(Q_3),size=(P_3+1,Q_3)) #
  initialise weights for layer W_3

```

Listing 4: Implementation of the multi-layer architecture

Forward pass

In a multi-layer architecture, the forward pass involves sequentially passing the input through each layer of the network. In this implementation, the data propagates through three weight matrices (W_1 , W_2 , and W_3). Since the weight matrices are designed to include a bias term, a constant value of 1 is added to the activation vector v_i using a vertical stacking operation to form the input vector u_{i+1} for the following layer. Listing 5 below demonstrates the implementation of the forward pass for the multi-layer architecture. The forward pass follows a chain: the features u_1 produce activations v_1 in the first hidden layer, which are then used to compute the activations v_2 in the second hidden layer, and finally the output y_{hat} is generated from the activations of the last layer.

```

1 u_1 = np.hstack([np.ones((1,1)),train_X[i:(i+1),:]]) .T # extract i'th
  input from training set and add bias
2 z_1,v_1 = forward(u_1,W_1) # calculate z_i and v_i for layer 1
3 u_2 = np.vstack([np.ones((1, 1)), v_1]) # Add bias to the input of the
  second layer
4 z_2, v_2 = forward(u_2, W_2) # Calculate z_i and v_i for layer 2
5 u_3 = np.vstack([np.ones((1,1)), v_2]) # Add bias to the input of the
  third layer
6 z_3, v_3 = forward(u_3, W_3) # Calculate z_i and v_i for layer 3
7 y_hat = v_3 # Final prediction is the activation of the last layer

```

Listing 5: Implementation of the forward pass for the multi-layer architecture

Backward pass

The backpropagation process involves calculating the gradients for each layer, starting from the output layer and moving backward through the network. In this implementation, the error signal Λ is initialised using the derivative of the binary cross-entropy loss and is iteratively updated at each layer. This iterative process ensures that each weight matrix receives the appropriate gradient information as described in the following equation.

$$\frac{\partial J}{\partial W_i} = \underbrace{\frac{\partial z_i}{\partial W_i}}_{u_i} \cdot \underbrace{\frac{\partial v_i}{\partial z_i}}_{\sigma'(z_i)} \cdot \Lambda \quad (4)$$

The gradients for each layer are accumulated in the `avg_grads` list, which is used to update the weights after processing a batch of data. The weights are updated using a gradient descent step as

described by the following formula:

$$W_i = W_i - \alpha \cdot \frac{\partial J}{\partial W_i} \quad (5)$$

In this equation, α represents the learning rate. Listing 6 below demonstrates the implementation of the backward pass for the multi-layer architecture.

```

1 Lambda = loss_derivative(y,y_hat) # initialise Lambda
2 dJ_dw3, Lambda = backward(u_3, v_3, W_3, Lambda) # Calculate derivatives
   for layer 3 by backpropagation
3 avg_grads[2] += dJ_dw3 # accumulate gradients for W_3
4 dJ_dw2, Lambda = backward(u_2, v_2, W_2, Lambda) # Calculate derivatives
   for layer 2 by backpropagation
5 avg_grads[1] += dJ_dw2 # accumulate gradients for W_2
6 dJ_dw1, Lambda = backward(u_1, v_1, W_1, Lambda) # calculate derivatives by
   backpropagation
7 avg_grads[0] += dJ_dw1 # accumulate gradients for W_1

```

Listing 6: Implementation of the backward pass for the multi-layer architecture

Decision surface

After training the multi-layer perceptron (with a learning rate of 0.1, 100 epochs and a batch size of 1) on the `two_blob` dataset, test set loss is analysed and the decision surface is visualised. Based on the training output, the model performance evolved as follows:

1. Epoch 0: test loss of -0.6572
2. Epoch 99: test loss of -0.2042

Table 2 illustrates model output over the last 10 epochs of training, whereas the plot, shown in Figure 4, displays the non-linear decision boundary that effectively separates the two classes.

Table 2: Training and test loss for the final 10 epochs of the single-layer perceptron on the `two_blob` dataset.

Epoch	Training Loss	Test Loss
90	-0.1618	-0.2166
91	-0.1545	-0.1885
92	-0.1436	-0.2212
93	-0.1369	-0.2294
94	-0.1278	-0.2491
95	-0.1302	-0.1579
96	-0.1336	-0.2296
97	-0.1310	-0.1639
98	-0.1227	-0.1945
99	-0.1153	-0.2042




Figure 4: Decision surface of the multi-layer perceptron on the `two_blob` dataset, demonstrating a non-linear decision boundary.

The above results demonstrate that the multi-layer architecture has successfully learned to capture the complex patterns in the data, which a single-layer perceptron would not have been able to achieve.

Conclusion

This practical successfully implemented both a single-layer and a multi-layer perceptron from fundamental principles to classify two different datasets. The following key conclusions can be made:

- The single-layer perceptron was sufficient to achieve optimal performance on the linearly separable `simple` dataset, as evidenced by the significant reduction in loss and the linear decision surface.
- The multi-layer perceptron was necessary to capture the non-linear patterns in the `two_blob` dataset, as demonstrated by the non-linear decision boundary and the reduction in loss over training.

Ultimately, this practical highlights the power of stacking single linear perceptron units to create complex architectures capable of capturing highly complex classification boundaries.